

```

import pandas as pd
import numpy as np
from scipy import stats
from sklearn.preprocessing import MinMaxScaler, OneHotEncoder
from sklearn.feature_selection import VarianceThreshold

df = pd.read_csv('/content/Cars_India_dataset.csv')

df.head()

{"type": "dataframe", "variable_name": "df"}

```

STEP-2:

DATA CLEANING AND PRE-PROCESSING TASK

```

import pandas as pd

# Data Cleaning Tasks

# 1. Handling Missing Values
# Check for missing values
missing_values = df.isnull().sum()
print("Missing Values:")
print(missing_values)

# Drop rows with missing values in specified columns
columns_to_check = ['Displacement cc', 'Emission Type', 'Fuel Tank Capacity', 'Transmission', 'No_of_Cylinders', 'Maker']
df.dropna(subset=columns_to_check, inplace=True)

# Drop the specified columns
columns_to_drop = ['Turning Radius', 'Boot Space', 'Model']
df.drop(columns=columns_to_drop, inplace=True)

# Confirm the changes
print(df.head())
missing_values = df.isnull().sum()

print(missing_values)

```

| Missing Values: | |
|-----------------|---|
| ID | 0 |
| Model | 0 |
| Maker | 0 |
| Type | 0 |
| Seats | 0 |
| Displacement cc | 1 |
| Length mm | 0 |
| Width | 0 |

| | |
|--------------------|----|
| Height | 0 |
| Wheelbase | 0 |
| No_of_Cylinders | 1 |
| Fuel | 0 |
| Engine Type | 0 |
| Transmission | 1 |
| Front Brake | 0 |
| Rear Brake | 0 |
| Drive | 0 |
| Turning Radius | 19 |
| Fuel Tank Capacity | 1 |
| Boot Space | 15 |
| Fuel Efficiency | 0 |
| Emission Type | 1 |
| Tyre Size | 0 |
| Variants | 0 |
| NCAP Rating | 0 |

dtype: int64

| ID | Maker | Type | Seats | Displacement cc | Length mm | |
|----|-------|------------|---------------|-----------------|-----------|------|
| 0 | 24 | Nissan | Sports Car | 2 | 3799.0 | 4710 |
| 1 | 40 | Mahindra | SUV | 6 | 2000.0 | 4662 |
| 2 | 2 | Volkswagen | Sedan | 5 | 1498.0 | 4561 |
| 3 | 9 | Honda | Compact Sedan | 5 | 1199.0 | 3995 |
| 4 | 8 | Honda | Compact Sedan | 5 | 1199.0 | 3995 |

| Height | Wheelbase | No_of_Cylinders | ... | Transmission | Front Brake |
|--------|-----------|-----------------|-----|--------------|-------------|
| 0 | 1370 | 2780 | 6.0 | 6-Speed DCT | Disc |
| 1 | 1857 | 2750 | 4.0 | 6 MT | Disc |
| 2 | 1507 | 2651 | 4.0 | 7-Speed DSG | Disc |
| 3 | 1501 | 2470 | 4.0 | CVT | Disc |
| 4 | 1501 | 2470 | 4.0 | 5 MT | Disc |

| Rear Brake | Drive | Fuel Tank Capacity | Fuel Efficiency | Emission |
|------------|-----------|--------------------|-----------------|----------|
| 0 | Disc 4 WD | 74.0 | 10.16 | Euro 6 |
| 1 | Disc 2WD | 57.0 | 18.60 | BS VI |
| 2 | Drum 2WD | 45.0 | 18.67 | BS VI |
| 3 | Drum 2WD | 35.0 | 18.30 | BS VI |
| 4 | Drum 2WD | 35.0 | 18.60 | BS VI |

| | Tyre Size | Variants | NCAP Rating |
|---|------------|----------|-------------|
| 0 | 285/35/20 | 1 | Not Tested |
| 1 | 255/16 R18 | 1 | Not Rated |
| 2 | 205/55 R16 | 1 | Not Tested |
| 3 | 175/65 R15 | 2 | 4 |
| 4 | 175/65 R15 | 3 | 4 |

[5 rows x 22 columns]

```
ID 0
Maker 0
Type 0
Seats 0
Displacement cc 1
Length mm 0
Width 0
Height 0
Wheelbase 0
No_of_Cylinders 1
Fuel 0
Engine Type 0
Transmission 1
Front Brake 0
Rear Brake 0
Drive 0
Fuel Tank Capacity 1
Fuel Efficiency 0
Emission Type 1
Tyre Size 0
Variants 0
NCAP Rating 0
dtype: int64
```

```
df.head()
```

```
{"type": "dataframe", "variable_name": "df"}
```

```
data = df
```

```
# 2. Checking for Duplicates
```

```
duplicate_rows = data.duplicated()
```

```
print("\nDuplicate Rows:")
```

```
print(df[duplicate_rows])
```

```
# df = data.drop_duplicates()
```

Duplicate Rows:

Empty DataFrame

Columns: [ID, Maker, Type, Seats, Displacement cc, Length mm, Width, Height, Wheelbase, No_of_Cylinders, Fuel, Engine Type, Transmission, Front Brake, Rear Brake, Drive, Fuel Tank Capacity, Fuel Efficiency,

```

Emission Type, Tyre Size, Variants, NCAP Rating]
Index: []

[0 rows x 22 columns]

# Data Preprocessing Tasks
# 1. Normalization/Scaling
from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()
data[['Length mm', 'Width', 'Height']] =
scaler.fit_transform(data[['Length mm', 'Width', 'Height']])

data.head()

{"type": "dataframe", "variable_name": "data"}

# 3. One-Hot Encoding
data = pd.get_dummies(data, columns=['Fuel', 'Type', 'Engine
Type', 'Transmission', 'Front Brake', 'Rear Brake', 'Drive', 'Tyre
Size', 'NCAP Rating', "Emission Type"])

```

STEP - 3

REGRESSION

```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, classification_report
from sklearn.impute import SimpleImputer

# Step 3: Impute missing values
imputer = SimpleImputer()
X_imputed = imputer.fit_transform(X)

# Step 4: Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_imputed, y,
test_size=0.2, random_state=42)

# Step 5: Initialize the Decision Tree Classifier
model = DecisionTreeClassifier(random_state=42)

# Step 6: Train the model on the training data
model.fit(X_train, y_train)

y_pred = model.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

```

```
# Classification report
print("\nClassification Report:")
print(classification_report(y_test, y_pred))
```

```
ID          0
Maker        0
Seats        0
Displacement cc  1
Length mm    0
..
NCAP Rating_6  0
NCAP Rating_Not Rated  0
NCAP Rating_Not Tested  0
Emission Type_BS VI  0
Emission Type_Euro 6  0
Length: 150, dtype: int64
Accuracy: 0.90625
```

Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| Citroen | 0.00 | 0.00 | 0.00 | 1 |
| Honda | 1.00 | 0.50 | 0.67 | 2 |
| Hyundai | 0.91 | 0.91 | 0.91 | 11 |
| Kia | 1.00 | 1.00 | 1.00 | 6 |
| Mahindra | 0.75 | 1.00 | 0.86 | 3 |
| Nissan | 1.00 | 1.00 | 1.00 | 2 |
| Renault | 1.00 | 1.00 | 1.00 | 1 |
| Tata | 1.00 | 1.00 | 1.00 | 2 |
| Toyota | 1.00 | 1.00 | 1.00 | 3 |
| Volkswagen | 0.50 | 1.00 | 0.67 | 1 |
| accuracy | | | 0.91 | 32 |
| macro avg | 0.82 | 0.84 | 0.81 | 32 |
| weighted avg | 0.90 | 0.91 | 0.89 | 32 |

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use
```

```
`zero_division` parameter to control this behavior.  
_warn_prf(average, modifier, msg_start, len(result))
```

1.Outlier Detection

```
from scipy.stats import zscore
import pandas as pd
import numpy as np

df=pd.read_csv('Car Sales.xlsx - car_data.csv')
df.head()

# Select numerical features, excluding the target column
numerical_features = df.select_dtypes(include=np.number).columns[:-1]

# Calculate Z-scores for numerical features only
z_scores = np.abs(zscore(df[numerical_features]))

# Define threshold
threshold = 3

# Detect outliers
outliers_zscore = (z_scores > threshold).any(axis=1)

# Display outliers
print("Z-Score Method Outliers:\n", df[outliers_zscore])
```

➡ Z-Score Method Outliers:

| | Car_id | Date | Customer Name | Gender | Annual Income | \ |
|-------|--------------|------------|---------------|--------|---------------|---|
| 8 | C_CND_000009 | 1/2/2022 | Naomi | Male | 815000 | |
| 48 | C_CND_000049 | 1/3/2022 | Valentine | Male | 406000 | |
| 122 | C_CND_000123 | 1/9/2022 | Benott | Male | 300000 | |
| 170 | C_CND_000171 | 1/13/2022 | Jordan | Male | 13500 | |
| 269 | C_CND_000270 | 1/27/2022 | Destiny | Female | 825000 | |
| ... | ... | ... | ... | ... | ... | |
| 23711 | C_CND_023712 | 12/29/2023 | Gilian | Male | 900000 | |
| 23762 | C_CND_023763 | 12/29/2023 | Loan | Male | 1234500 | |
| 23769 | C_CND_023770 | 12/29/2023 | Shyrel | Male | 4111000 | |
| 23809 | C_CND_023810 | 12/29/2023 | Louen | Male | 2065000 | |
| 23839 | C_CND_023840 | 12/30/2023 | Sofia | Female | 555000 | |

| | Dealer Name | Company | \ |
|-------|--|------------|---|
| 8 | Rabun Used Car Sales | Chevrolet | |
| 48 | Race Car Help | Nissan | |
| 122 | Clay Johnson Auto Sales | Cadillac | |
| 170 | Scrivener Performance Engineering | Lincoln | |
| 269 | Pars Auto Sales | Toyota | |
| ... | ... | ... | |
| 23711 | Capitol KIA | Cadillac | |
| 23762 | Progressive Shippers Cooperative Association | No Lincoln | |
| 23769 | Suburban Ford | Volkswagen | |
| 23809 | U-Haul CO | Lincoln | |
| 23839 | Suburban Ford | Cadillac | |

| | Model | Engine | Transmission | Color | \ |
|-------|-------------|---------------------------|--------------|------------|---|
| 8 | Malibu | Overhead Camshaft | Manual | Pale White | |
| 48 | Altima | DoubleÃ Overhead Camshaft | Auto | Pale White | |
| 122 | Eldorado | DoubleÃ Overhead Camshaft | Auto | Pale White | |
| 170 | Continental | Overhead Camshaft | Manual | Pale White | |
| 269 | Tacoma | Overhead Camshaft | Manual | Pale White | |
| ... | ... | ... | ... | ... | |
| 23711 | Eldorado | Overhead Camshaft | Manual | Black | |
| 23762 | Continental | Overhead Camshaft | Manual | Red | |
| 23769 | GTI | DoubleÃ Overhead Camshaft | Auto | Pale White | |
| 23809 | Continental | Overhead Camshaft | Manual | Red | |
| 23839 | Catera | DoubleÃ Overhead Camshaft | Auto | Pale White | |

| | Price (\$) | Dealer_No | Body Style | Phone | Dealer_Region |
|-------|------------|------------|------------|---------|---------------|
| 8 | 82000 | 85257-3102 | Hardtop | 7194857 | Pasco |
| 48 | 20000 | 78758-7841 | Hatchback | 7117432 | Austin |
| 122 | 31000 | 78758-7841 | Passenger | 8668755 | Austin |
| 170 | 82000 | 38701-8047 | Passenger | 6642461 | Greenville |
| 269 | 82000 | 38701-8047 | Hatchback | 7848361 | Greenville |
| ... | ... | ... | ... | ... | ... |
| 23711 | 85000 | 38701-8047 | Passenger | 7788669 | Greenville |
| 23762 | 82450 | 53546-9427 | Passenger | 7468114 | Austin |
| 23769 | 20100 | 53546-9427 | SUV | 8363552 | Austin |
| 23809 | 82500 | 78758-7841 | Passenger | 6406323 | Aurora |
| 23839 | 75000 | 53546-9427 | Hatchback | 7752902 | Janesville |

[680 rows x 16 columns]

Double-click (or enter) to edit

```
from sklearn.ensemble import IsolationForest
```

```
# Initialize Isolation Forest
iso_forest = IsolationForest(contamination=0.1, random_state=42)

# Select numerical features, excluding the target column and 'Car_id'
numerical_features = df.select_dtypes(include=np.number).columns.difference(['outlier_iso'])

# Fit and predict using only numerical features
df['outlier_iso'] = iso_forest.fit_predict(df[numerical_features])

# Display outliers
print("Isolation Forest Outliers:\n", df[df['outlier_iso'] == -1])
```

```
↗ Isolation Forest Outliers:
```

| | Car_id | Date | Customer Name | Gender | Annual Income | \ |
|-------|--------------|------------|---------------|--------|---------------|---|
| 8 | C_CND_000009 | 1/2/2022 | Naomi | Male | 815000 | |
| 21 | C_CND_000022 | 1/2/2022 | Joshua | Male | 2500000 | |
| 37 | C_CND_000038 | 1/3/2022 | Haylee | Male | 13500 | |
| 48 | C_CND_000049 | 1/3/2022 | Valentine | Male | 4060000 | |
| 65 | C_CND_000066 | 1/4/2022 | Annaelle | Male | 1650000 | |
| ... | ... | ... | ... | ... | ... | |
| 23885 | C_CND_023886 | 12/31/2023 | Jeremias | Male | 1890000 | |
| 23890 | C_CND_023891 | 12/31/2023 | Joaquin | Male | 2450000 | |
| 23891 | C_CND_023892 | 12/31/2023 | Annabelle | Male | 2340000 | |
| 23895 | C_CND_023896 | 12/31/2023 | Sima | Male | 965000 | |
| 23899 | C_CND_023900 | 12/31/2023 | Yuna | Male | 13500 | |

| | Dealer Name | Company | \ |
|-------|---|-----------|---|
| 8 | Rabun Used Car Sales | Chevrolet | |
| 21 | Classic Chevy | Infiniti | |
| 37 | Gartner Buick Hyundai Saab | Buick | |
| 48 | Race Car Help | Nissan | |
| 65 | Star Enterprises Inc | Buick | |
| ... | ... | ... | |
| 23885 | Progressive Shippers Cooperative Association No | Ford | |
| 23890 | Saab-Belle Dodge | Dodge | |
| 23891 | Ryder Truck Rental and Leasing | Chevrolet | |
| 23895 | Progressive Shippers Cooperative Association No | Mercury | |
| 23899 | U-Haul CO | Buick | |

| | Model | Engine | Transmission | Color | \ |
|-------|-------------|---------------------------|--------------|------------|---|
| 8 | Malibu | Overhead Camshaft | Manual | Pale White | |
| 21 | I30 | DoubleÃ Overhead Camshaft | Auto | Black | |
| 37 | Park Avenue | DoubleÃ Overhead Camshaft | Auto | Black | |
| 48 | Altima | DoubleÃ Overhead Camshaft | Auto | Pale White | |
| 65 | Park Avenue | DoubleÃ Overhead Camshaft | Auto | Black | |
| ... | ... | ... | ... | ... | |
| 23885 | Ranger | Overhead Camshaft | Manual | Black | |
| 23890 | Ram Pickup | Overhead Camshaft | Manual | Pale White | |
| 23891 | Corvette | DoubleÃ Overhead Camshaft | Auto | Pale White | |
| 23895 | Sable | Overhead Camshaft | Manual | Red | |
| 23899 | Park Avenue | DoubleÃ Overhead Camshaft | Auto | Pale White | |

| | Price (\$) | Dealer_No | Body Style | Phone | Dealer_Region | outlier_iso |
|-------|------------|------------|------------|---------|---------------|-------------|
| 8 | 82000 | 85257-3102 | Hardtop | 7194857 | Pasco | -1 |
| 21 | 21000 | 85257-3102 | Hardtop | 6183219 | Austin | -1 |
| 37 | 61000 | 38701-8047 | Hatchback | 7438037 | Greenville | -1 |
| 48 | 20000 | 78758-7841 | Hatchback | 7117432 | Austin | -1 |
| 65 | 61000 | 99301-3882 | Hatchback | 8380613 | Pasco | -1 |
| ... | ... | ... | ... | ... | ... | ... |
| 23885 | 18000 | 53546-9427 | Hardtop | 6009530 | Janesville | -1 |
| 23890 | 20001 | 60504-7114 | Hardtop | 6172324 | Aurora | -1 |
| 23891 | 46000 | 06457-3834 | SUV | 6435802 | Middletown | -1 |
| 23895 | 61000 | 53546-9427 | Sedan | 8439821 | Middletown | -1 |
| 23899 | 62000 | 78758-7841 | Hatchback | 8384785 | Aurora | -1 |

[2391 rows x 17 columns]

```
from sklearn.neighbors import LocalOutlierFactor
import pandas as pd

# Initialize Local Outlier Factor
lof = LocalOutlierFactor(n_neighbors=20, contamination=0.1)

# Select only numerical features for outlier detection
numerical_features = df.select_dtypes(include=['number']) # Select numerical columns

# Fit and predict on numerical features only
outliers_lof = lof.fit_predict(numerical_features)

# Add LOF results to the dataframe
df['outlier_lof'] = outliers_lof

# Display outliers
print("Local Outlier Factor Outliers:\n", df[df['outlier_lof'] == -1])
```



```

Local Outlier Factor Outliers:
      Car_id      Date Customer Name Gender Annual Income \
7      C_CND_000008  1/2/2022      Graham Male      13500
8      C_CND_000009  1/2/2022      Naomi Male      815000
11     C_CND_000012  1/2/2022      Amar'E Male      13500
28     C_CND_000029  1/2/2022      Sloane Male      13500
30     C_CND_000031  1/2/2022      Sophia Male     210000
...
23839  C_CND_023840  12/30/2023      Sofia Female    555000
23846  C_CND_023847  12/30/2023      Sylvia Female    925000
23873  C_CND_023874  12/31/2023      Gabriel Male      13500
23881  C_CND_023882  12/31/2023      Vicky Male      843000
23899  C_CND_023900  12/31/2023      Yuna Male      13500

      Dealer_Name      Company      Model \
7      U-Haul CO      Mitsubishi Galant
8      Rabun Used Car Sales Chevrolet Malibu
11     Race Car Help      Nissan Pathfinder
28     Race Car Help      Chrysler LHS
30     Saab-Belle Dodge Mitsubishi 3000GT
...
23839  Suburban Ford      Cadillac Catera
23846  Race Car Help      Oldsmobile Aurora
23873  Saab-Belle Dodge      Subaru Outback
23881  Star Enterprises Inc Lexus LS400
23899  U-Haul CO      Buick Park Avenue

      Engine Transmission      Color Price ($) \
7      DoubleÃ Overhead Camshaft Auto Pale White 42000
8      Overhead Camshaft Manual Pale White 82000
11     DoubleÃ Overhead Camshaft Auto Pale White 46000
28     Overhead Camshaft Manual Pale White 41000
30     Overhead Camshaft Manual Pale White 20000
...
23839  DoubleÃ Overhead Camshaft Auto Pale White 75000
23846  Overhead Camshaft Manual Red 71000
23873  Overhead Camshaft Manual Red 49000
23881  Overhead Camshaft Manual Black 69001
23899  DoubleÃ Overhead Camshaft Auto Pale White 62000

      Dealer_No Body Style Phone Dealer_Region outlier_iso outlier_lof
7      78758-7841 Passenger 6206512 Austin 1 -1
8      85257-3102 Hardtop 7194857 Pasco -1 -1
11     78758-7841 Hardtop 7288103 Pasco 1 -1
28     78758-7841 Hatchback 6292720 Janesville 1 -1
30     60504-7114 Sedan 8847858 Austin 1 -1
...
23839  53546-9427 Hatchback 7752902 Janesville -1 -1
23846  78758-7841 Passenger 7265067 Austin -1 -1
23873  60504-7114 SUV 7090003 Aurora 1 -1
23881  99301-3882 Sedan 7011127 Pasco -1 -1
23899  78758-7841 Hatchback 8384785 Aurora -1 -1

```

[2391 rows x 18 columns]

```

import seaborn as sns
import matplotlib.pyplot as plt

# Visualize each feature using boxplots
# Calculate the number of rows and columns for subplots dynamically
num_cols = len(df.columns[:-1]) # Number of features to plot
num_rows = (num_cols + 1) // 2 # Calculate rows, ensuring enough space

plt.figure(figsize=(12, 8))
for i, column in enumerate(df.columns[:-1]): # Exclude the target and outlier columns
    plt.subplot(num_rows, 2, i + 1) # Use calculated rows and columns
    sns.boxplot(y=df[column], color="lightblue")
    plt.title(f'Boxplot of {column}')

plt.tight_layout()
plt.show()

...

```

2. Text Mining

```
import pandas as pd
import nltk
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from collections import Counter

# Download necessary NLTK data
nltk.download('punkt')
nltk.download('stopwords')
# Download the 'punkt_tab' data package for sentence tokenization
nltk.download('punkt_tab')

# Load dataset (replace 'your_dataset.csv' with your actual dataset)
df = pd.read_csv('Car Sales.xlsx - car_data.csv')

# Assuming 'text_column' is the column containing text data
text_data = df['Dealer_Region'].tolist()

# Tokenize and preprocess text
stop_words = set(stopwords.words('english'))

def preprocess_text(text):
    tokens = word_tokenize(text.lower())
    return [token for token in tokens if token.isalnum() and token not in stop_words]

preprocessed_data = [preprocess_text(text) for text in text_data]

# Count word frequencies
all_words = [word for text in preprocessed_data for word in text]
word_freq = Counter(all_words)

# Print top 10 most common words
print("Top 10 most common words:")
for word, count in word_freq.most_common(10):
    print(f"{word}: {count}")


# Basic sentiment analysis (you may need to install TextBlob)
from textblob import TextBlob

def get_sentiment(text):
    return TextBlob(text).sentiment.polarity

# Assuming 'text_column' is the correct column name
# Replace 'text_column' with the actual column name if it's different
df['sentiment'] = df['Dealer_Region'].apply(get_sentiment)

print("\nAverage sentiment score:", df['sentiment'].mean())

# You can add more advanced text mining techniques here, such as:
# - Topic modeling (e.g., using Latent Dirichlet Allocation)
# - Named Entity Recognition
# - Text classification
# - Word embeddings (e.g., Word2Vec, GloVe)
```

 [nltk_data] Downloading package punkt to /root/nltk_data...
 [nltk_data] Package punkt is already up-to-date!
 [nltk_data] Downloading package stopwords to /root/nltk_data...
 [nltk_data] Package stopwords is already up-to-date!
 [nltk_data] Downloading package punkt_tab to /root/nltk_data...
 [nltk_data] Package punkt_tab is already up-to-date!
 Top 10 most common words:
 austin: 4135
 janesville: 3821
 scottsdale: 3433
 pasco: 3131
 aurora: 3130
 middletown: 3128
 greenville: 3128
 Average sentiment score: 0.0


```
# import python libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# import csv file
df = pd.read_csv('/content/drive/MyDrive/Colab Notebooks/DAV
Project/Diwali Sales Data.csv', encoding= 'unicode_escape')

from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
df.shape
```

```
(11251, 15)
```

```
df.head(20)
```

| | User_ID | Cust_name | Product_ID | Gender | Age Group | Age |
|----|---------|-----------|------------|--------|-----------|-----|
| 0 | 1002903 | Sanskriti | P00125942 | F | 26-35 | 28 |
| 1 | 1000732 | Kartik | P00110942 | F | 26-35 | 35 |
| 2 | 1001990 | Bindu | P00118542 | F | 26-35 | 35 |
| 3 | 1001425 | Sudevi | P00237842 | M | 0-17 | 16 |
| 4 | 1000588 | Joni | P00057942 | M | 26-35 | 28 |
| 5 | 1000588 | Joni | P00057942 | M | 26-35 | 28 |
| 6 | 1001132 | Balk | P00018042 | F | 18-25 | 25 |
| 7 | 1002092 | Shivangi | P00273442 | F | 55+ | 61 |
| 8 | 1003224 | Kushal | P00205642 | M | 26-35 | 35 |
| 9 | 1003650 | Ginny | P00031142 | F | 26-35 | 26 |
| 10 | 1003829 | Harshita | P00200842 | M | 26-35 | 34 |
| 11 | 1000214 | Kargatis | P00119142 | F | 18-25 | 20 |
| 12 | 1004035 | Elijah | P00080342 | F | 18-25 | 20 |
| 13 | 1001680 | Vasudev | P00324942 | M | 26-35 | 26 |

| | | | | | | |
|----|---------|---------|-----------|---|-------|----|
| 14 | 1003858 | Cano | P00293742 | M | 46-50 | 46 |
| 1 | | | | | | |
| 15 | 1000813 | Lauren | P00289942 | F | 18-25 | 24 |
| 0 | | | | | | |
| 16 | 1005447 | Amy | P00275642 | F | 46-50 | 48 |
| 1 | | | | | | |
| 17 | 1001193 | Mick | P00004842 | F | 26-35 | 29 |
| 0 | | | | | | |
| 18 | 1001883 | Praneet | P00029842 | M | 51-55 | 54 |
| 1 | | | | | | |
| 19 | 1001883 | Praneet | P00029842 | M | 51-55 | 54 |
| 1 | | | | | | |

| Orders \ | State | Zone | Occupation | Product_Category |
|----------|------------------|----------|-----------------|------------------|
| 0 | Maharashtra | Western | Healthcare | Auto |
| 1 | | | | |
| 1 | Andhra Pradesh | Southern | Govt | Auto |
| 3 | | | | |
| 2 | Uttar Pradesh | Central | Automobile | Auto |
| 3 | | | | |
| 3 | Karnataka | Southern | Construction | Auto |
| 2 | | | | |
| 4 | Gujarat | Western | Food Processing | Auto |
| 2 | | | | |
| 5 | Himachal Pradesh | Northern | Food Processing | Auto |
| 1 | | | | |
| 6 | Uttar Pradesh | Central | Lawyer | Auto |
| 4 | | | | |
| 7 | Maharashtra | Western | IT Sector | Auto |
| 1 | | | | |
| 8 | Uttar Pradesh | Central | Govt | Auto |
| 2 | | | | |
| 9 | Andhra Pradesh | Southern | Media | Auto |
| 4 | | | | |
| 10 | Delhi | Central | Banking | Auto |
| 1 | | | | |
| 11 | Andhra Pradesh | Southern | Retail | Auto |
| 2 | | | | |
| 12 | Andhra Pradesh | Southern | IT Sector | Auto |
| 2 | | | | |
| 13 | Andhra Pradesh | Southern | Automobile | Auto |
| 4 | | | | |
| 14 | Madhya Pradesh | Central | Hospitality | Auto |
| 3 | | | | |
| 15 | Andhra Pradesh | Southern | Govt | Auto |
| 2 | | | | |
| 16 | Andhra Pradesh | Southern | IT Sector | Auto |
| 3 | | | | |

| | | | | |
|----|----------------|----------|-------------|------|
| 17 | Andhra Pradesh | Southern | Aviation | Auto |
| 1 | | | | |
| 18 | Uttar Pradesh | Central | Hospitality | Auto |
| 1 | | | | |
| 19 | Uttar Pradesh | Central | Hospitality | Auto |
| 1 | | | | |

| | Amount | Status | unnamed1 |
|----|----------|--------|----------|
| 0 | 23952.00 | NaN | NaN |
| 1 | 23934.00 | NaN | NaN |
| 2 | 23924.00 | NaN | NaN |
| 3 | 23912.00 | NaN | NaN |
| 4 | 23877.00 | NaN | NaN |
| 5 | 23877.00 | NaN | NaN |
| 6 | 23841.00 | NaN | NaN |
| 7 | NaN | NaN | NaN |
| 8 | 23809.00 | NaN | NaN |
| 9 | 23799.99 | NaN | NaN |
| 10 | 23770.00 | NaN | NaN |
| 11 | 23752.00 | NaN | NaN |
| 12 | 23730.00 | NaN | NaN |
| 13 | 23718.00 | NaN | NaN |
| 14 | NaN | NaN | NaN |
| 15 | 23664.00 | NaN | NaN |
| 16 | NaN | NaN | NaN |
| 17 | 23619.00 | NaN | NaN |
| 18 | 23568.00 | NaN | NaN |
| 19 | 23568.00 | NaN | NaN |

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 11251 entries, 0 to 11250
Data columns (total 15 columns):
```

| # | Column | Non-Null Count | Dtype |
|-----|------------------|----------------|---------|
| --- | ----- | ----- | ----- |
| 0 | User_ID | 11251 non-null | int64 |
| 1 | Cust_name | 11251 non-null | object |
| 2 | Product_ID | 11251 non-null | object |
| 3 | Gender | 11251 non-null | object |
| 4 | Age Group | 11251 non-null | object |
| 5 | Age | 11251 non-null | int64 |
| 6 | Marital_Status | 11251 non-null | int64 |
| 7 | State | 11251 non-null | object |
| 8 | Zone | 11251 non-null | object |
| 9 | Occupation | 11251 non-null | object |
| 10 | Product_Category | 11251 non-null | object |
| 11 | Orders | 11251 non-null | int64 |
| 12 | Amount | 11239 non-null | float64 |
| 13 | Status | 0 non-null | float64 |

```
14  unnamed1          0 non-null      float64
dtypes: float64(3), int64(4), object(8)
memory usage: 1.3+ MB
```

```
#drop unrelated/blank columns
```

```
df.drop(['Status', 'unnamed1'], axis=1, inplace=True)
```

```
#check for null values
```

```
pd.isnull(df).sum()
```

```
User_ID          0
Cust_name        0
Product_ID       0
Gender           0
Age Group        0
Age              0
Marital_Status   0
State            0
Zone             0
Occupation       0
Product_Category 0
Orders           0
Amount           12
dtype: int64
```

```
# drop null values
```

```
df.dropna(inplace=True)
```

```
# change data type
```

```
df['Amount'] = df['Amount'].astype('int')
```

```
df['Amount'].dtypes
```

```
dtype('int64')
```

```
df.columns
```

```
Index(['User_ID', 'Cust_name', 'Product_ID', 'Gender', 'Age Group',
      'Age',
      'Marital_Status', 'State', 'Zone', 'Occupation',
      'Product_Category',
      'Orders', 'Amount'],
      dtype='object')
```

```
df
```

```
# describe() method returns description of the data in the DataFrame
(i.e. count, mean, std, etc)
```

```
df.describe()
```

```
      User_ID      Age  Marital_Status      Orders
Amount \
```

| | | | | |
|--------------|--------------|--------------|--------------|--------------|
| count | 1.125100e+04 | 11251.000000 | 11251.000000 | 11251.000000 |
| 11239.000000 | | | | |
| mean | 1.003004e+06 | 35.421207 | 0.420318 | 2.489290 |
| 9453.610858 | | | | |
| std | 1.716125e+03 | 12.754122 | 0.493632 | 1.115047 |
| 5222.355869 | | | | |
| min | 1.000001e+06 | 12.000000 | 0.000000 | 1.000000 |
| 188.000000 | | | | |
| 25% | 1.001492e+06 | 27.000000 | 0.000000 | 1.500000 |
| 5443.000000 | | | | |
| 50% | 1.003065e+06 | 33.000000 | 0.000000 | 2.000000 |
| 8109.000000 | | | | |
| 75% | 1.004430e+06 | 43.000000 | 1.000000 | 3.000000 |
| 12675.000000 | | | | |
| max | 1.006040e+06 | 92.000000 | 1.000000 | 4.000000 |
| 23952.000000 | | | | |

| | Status | unnamed1 |
|-------|--------|----------|
| count | 0.0 | 0.0 |
| mean | NaN | NaN |
| std | NaN | NaN |
| min | NaN | NaN |
| 25% | NaN | NaN |
| 50% | NaN | NaN |
| 75% | NaN | NaN |
| max | NaN | NaN |

```
# use describe() for specific columns
df[['Age', 'Orders', 'Amount']].describe()
```

| | Age | Orders | Amount |
|-------|--------------|--------------|--------------|
| count | 11251.000000 | 11251.000000 | 11239.000000 |
| mean | 35.421207 | 2.489290 | 9453.610858 |
| std | 12.754122 | 1.115047 | 5222.355869 |
| min | 12.000000 | 1.000000 | 188.000000 |
| 25% | 27.000000 | 1.500000 | 5443.000000 |
| 50% | 33.000000 | 2.000000 | 8109.000000 |
| 75% | 43.000000 | 3.000000 | 12675.000000 |
| max | 92.000000 | 4.000000 | 23952.000000 |

Exploratory Data Analysis


Gender

```
# plotting a bar chart for Gender and it's count

ax = sns.countplot(x = 'Gender', data = df)
```




```
for bars in ax.containers:  
    ax.bar_label(bars)
```



Ques1: What is the distribution of purchasing power among different genders, and how does it compare between males and females?

```
ax = sns.countplot(data = df, x = 'Age Group', hue = 'Gender')  
  
for bars in ax.containers:  
    ax.bar_label(bars)
```




From above graphs we can see that most of the buyers are females and even the purchasing power of females are greater than men

Age

Ques2: Most and least buyers are in which age group?

```
# Total Amount vs Age Group
sales_age = df.groupby(['Age Group'], as_index=False)
['Amount'].sum().sort_values(by='Amount', ascending=False)

sns.barplot(x = 'Age Group', y= 'Amount' ,data = sales_age)
<Axes: xlabel='Age Group', ylabel='Amount'>
```



From above graphs we can see that most of the buyers are of age group between 26-35 yrs and least are of age group between 0-17 yrs

State


Ques3: List the Top 10 States having most purchases

```
# total number of orders from top 10 states

sales_state = df.groupby(['State'], as_index=False)
['Orders'].sum().sort_values(by='Orders', ascending=False).head(10)

sns.set(rc={'figure.figsize':(15,5)})
sns.barplot(data = sales_state, x = 'State',y= 'Orders')

<Axes: xlabel='State', ylabel='Orders'>
```




Ques3: What are the key regions driving the highest number of orders and total sales?

```
# total amount/sales from top 10 states
```

```
sales_state = df.groupby(['State'], as_index=False)
['Amount'].sum().sort_values(by='Amount', ascending=False).head(10)
```

```
sns.set(rc={'figure.figsize':(15,5)})
sns.barplot(data = sales_state, x = 'State', y= 'Amount')
```

```
<Axes: xlabel='State', ylabel='Amount'>
```




From above graphs we can see that most of the orders & total sales/amount are from Uttar Pradesh, Maharashtra and Karnataka respectively

Marital Status

```
ax = sns.countplot(data = df, x = 'Marital_Status')
```

```
sns.set(rc={'figure.figsize':(7,5)})
```

```
for bars in ax.containers:
    ax.bar_label(bars)
```




Ques4: "What is the purchasing behavior and power of married women in the dataset?"

```
sales_state = df.groupby(['Marital_Status', 'Gender'], as_index=False)
['Amount'].sum().sort_values(by='Amount', ascending=False)

sns.set(rc={'figure.figsize':(6,5)})
sns.barplot(data = sales_state, x = 'Marital_Status', y= 'Amount',
hue='Gender')

<Axes: xlabel='Marital_Status', ylabel='Amount'>
```




From above graphs we can see that most of the buyers are married (women) and they have high purchasing power

Occupation

Ques5: "What is the occupational distribution of buyers in the dataset, and how does it vary across different sectors?"

```
sns.set(rc={'figure.figsize':(20,5)})
ax = sns.countplot(data = df, x = 'Occupation')


for bars in ax.containers:
    ax.bar_label(bars)
```



```
sales_state = df.groupby(['Occupation'], as_index=False)
['Amount'].sum().sort_values(by='Amount', ascending=False)

sns.set(rc={'figure.figsize':(20,5)})
sns.barplot(data = sales_state, x = 'Occupation',y= 'Amount')

<Axes: xlabel='Occupation', ylabel='Amount'>
```




From above graphs we can see that most of the buyers are working in IT, Healthcare and Aviation sector

Product Category

Ques 7: What is the distribution of sold products across various categories?


```
sns.set(rc={'figure.figsize':(20,5)})
ax = sns.countplot(data = df, x = 'Product_Category')

for bars in ax.containers:
    ax.bar_label(bars)
```



```
sales_state = df.groupby(['Product_Category'], as_index=False)
['Amount'].sum().sort_values(by='Amount', ascending=False).head(10)

sns.set(rc={'figure.figsize':(20,5)})
sns.barplot(data = sales_state, x = 'Product_Category',y= 'Amount')
<Axes: xlabel='Product_Category', ylabel='Amount'>
```



From above graphs we can see that most of the sold products are from Food, Clothing and Electronics category

Which Product ID have maximum number of orders?

```
sales_state = df.groupby(['Product_ID'], as_index=False)
['Orders'].sum().sort_values(by='Orders', ascending=False).head(10)

sns.set(rc={'figure.figsize':(20,5)})
sns.barplot(data = sales_state, x = 'Product_ID',y= 'Orders')
<Axes: xlabel='Product_ID', ylabel='Orders'>
```